

Classical Search Optimizations for Entity-Linking Systems

Raul E. Jordan

*A.B. Candidate in Computer Science.
Harvard University.*

RAULJORDAN@COLLEGE.HARVARD.EDU

Melissa C. Lee

*A.B. Candidate in Computer Science.
Harvard University*

LEE05@COLLEGE.HARVARD.EDU

Sameer N. Mehra

*A.B. Candidate in Computer Science.
Harvard University*

SMEHRA@COLLEGE.HARVARD.EDU

Abstract

This paper describes a classical search approach to linking entities in text to their corresponding wikipedia entries through optimizations of semantic-relatedness and a novel abstraction using existing methods. We will obtain a solution for a given input in this problem through the optimization of certain parameters and the confidence value of a keyword-link pair through a “relevance score” using methods such as Explicit Semantic Analysis. Most of this project will explore different methods of defining this relevance function and its results. We present an approach that has wide applicability to understanding natural language and making it easier for speech based agents to discover what a user is referring to based on the wikipedia knowledge base.

1. Introduction

Named-Entity Linking is a *explain*. This is a complex task that has yielded powerful results in recent years, and the strength of using a structured corpora of Wikipedia has given rise to a wide-variety of techniques in the fields of machine learning and artificial intelligence as part of the greater scope of natural language processing that seeks to identify entities effectively. On a narrower panorama, the task of entity linking based on wikipedia is a task we describe as assigning the most-related. This paper will focus on defining a robust measure of relatedness to achieve a classical search-based, artificial intelligence solution to the task of named-entity recognition to accomplishing effective entity linking by optimizing various methods of quantifying this most-related method. Being able to use a structured KB as a way of quickly identifying entities from context has a myriad of applications ranging from speech recognition, to intelligent personal assistants, which is especially relevant to recent systems such as Facebook’s “M”.

We will follow measures of semantic relatedness known from recent natural language processing research as ways of obtaining robust measures that determining whether or not an entity and link assignment is accurate enough.

1.1 A New Abstraction Based on Existing Methods

Our new approach formulates named-entity linking (NEL) as a search problem through a construction known as a relevance function. This is an abstraction that relies on existing algorithms, such as TF-IDF and Explicit Semantic Analysis (Hachey 40) to determine if a link is the most likely candidate for a given keyword. This function will map a keyword-link pairing to a score between 0 and 1 known as a *relevance score* based on the existing *semantic relatedness* algorithms such as the ones mentioned before. Our formulation abstracts most of the problem-specific details of NEL into simply searching a state space for the best possible pairing of entities to candidate links above a certain relevance score threshold. The key to our formulation is that every assignment of a keyword to a candidate link is crucial to determining the next possible assignment effectively. That is, every action is deterministic and there are multiple approaches our implementation follows to structure NEL as a problem that can be solved through algorithms such as Depth First Search or Breadth First Search. Much of our discussion focuses on the problem of assessing an effective relevance function mapping and taking into account previous link-keyword assignments to achieve a terminal state of our formulation.

1.2 Range of Problems Addressed

As a search problem, every state and action pair in a given search path will be important to the next state and action pair we determine. As previously mentioned, our solution will require that each previous link-keyword assignment provide crucial information for future assignments. One of the problems we encounter is how to determine the relevance between assignments as we traverse the state space of possible links a keyword can be paired to. In many cases, keywords in an input are highly related and one keyword can be immediately helpful in constraining the possible candidate state space of another keyword in the input. For example,

This is what we refer to as *linking-dependence*. However, it might not be the case that a particular future assignment has any dependence on the previous ones or the other entities around it. For example, assume that our input is “Steve Jobs was fantastic at Apple” and that we have successfully linked “Steve Jobs” to his wikipedia page with a high confidence (relevance score) value. Suppose now that we are assigning “Apple” to its relevant link and that its candidate link state space contains.

$$\text{CANDIDATES}(\textit{Apple}) = \{\textit{Apple}(\textit{Fruit}), \textit{AppleInc.}, \dots\}$$

However, taking into account that we have just assigned “Steve Jobs” before will restrict $\text{CANDIDATES}(\textit{Apple})$ to

$$\text{CANDIDATES}(\textit{Apple}) = \{\textit{Apple Inc.}\}$$

We will refer to this as linking-dependence. We need to talk more about the specific problems that arise with determining the most-constrained keyword effectively

2. Classical Search Formulation

Our system first takes in a block of text and preprocesses it to obtain the keywords and build up our initial state of our search problem as a dictionary where the keys are the

keywords and the values are their respective wikipedia page assignments which are initially set to $(None, 0)$ where the the second element of this tuple is a “relevance score” for each particular link-keyword pair.

For example, our initial input could be “An airplane has a heavy wing”, and our pre-processed initial state will end up being a dictionary of the form

$$\{\text{airplane} : (None, 0), \text{wing} : (None, 0)\}$$

This input is then passed into a search agent that uses optimal and suboptimal algorithms such as Breadth First Search, Depth First Search, Uniform Cost Search, and A^* Search with problem specific heuristics. Each action will involved choosing a keyword to assign and expanding the state space of candidate links the keyword can take on by taking into account previous keyword-link assignments. The solution returns a terminal state that has every word assigned a link with a relevance score above a certain threshold, γ based on the relevance function, such as

$$\{\text{airplane} : (\text{wikipedia.org/Airplane}, 0.98), \text{wing} : (\text{wikipedia.org/Wing}, 0.87)\}$$

There are additional considerations in our search formulation we have to take into account when formulating our state space this way. We will discuss these by going into detail of how a specific Algorithm such as Depth First Search in our system finds a highly relevant terminal state in the following section.

2.1 High Level Example: Depth First Search

We implement the Depth First Search algorithm specified in AIMA (citation needed) but with key differences in the how each node is expanded into all of its possible child nodes. Also, choosing a keyword to assign will be the key of a deterministic, classical search formulation and we will do so by taking into account previous keyword-link assignments. We will accomplish this through a method defined as the *most-constrained keyword heuristic*. The main driver of our implementation is obtaining the associated actions of these child nodes to append to the frontier of the search algorithm, which in this case is represented using a stack as a data structure. This is where relevance scores are computed and assigned to each possible keyword-link pairing.

Obtaining the child states for a given state will be as follows

1. Choose a keyword to assign based on the most-constrained keyword heuristic
2. Query the wikipedia API for every possible candidate link that keyword can take. Other literature on entity-linking mentions that candidate generation is critical for successful named-entity linking (Hachey 40), but our implementation already does this by abstracting this through the wikipedia API
3. Compute the relevance score between every possible link and the given keyword
4. Return a list of triples of the form

$$[(\text{keyword}, \text{link}, \text{score}), \dots]$$

that includes all possible candidate links and their associated score for a given keyword.

Afterwards, we build dictionaries for all of the possible resulting states from this list of successor states and we can push them to the search algorithms frontier.

Our algorithm ends when we find a state that has every keyword assigned with a relevance score of at least a certain value, γ we decide between 0 and 1. This logic is implemented in the `ISGOALSTATE` function which checks for a terminal state in our implementation as in the AIMA algorithmic description. This is a score cutoff that can be modified at run time and we will discuss its limitations and the benefits of choosing different values of γ in following sections. Now we discuss the aspect of our implementation that frames the problem as deterministic classical search and - the most-constrained keyword heuristic.

2.2 Most-Constrained Keyword

Given a state, how does the algorithm determine the keyword to assign based on the other assignments? Ideally, we want to restrict the state space of every unassigned keyword based on the other assignments to find the one that has the smallest state space of possible link as in the *Steve Jobs* and *Apple* example in section 1.2. This is necessary for our algorithm to make sense as a classical search problem. We want to be able to use current assignments to help determine future ones, and using this abstraction will allow us to implement our solution effectively. However, we run into the problem of linking-dependence vs. linking-independence as introduced in section 1.2. To solve this problem, we calculate the linking-dependence between two keywords by using the cosine similarity measure discussed in the following section. We then introduce an extra parameter, a *linking-dependence threshold* ψ , that can be modified at run-time arbitrarily. The efficacy of different choices of ψ will be discussed in section 4 as part of the algorithmic limitations of our implementation. If two keywords are dependent on each other enough to have a similarity score above ψ , the algorithm will restrict the state space of the keywords accordingly as in the example.

```
def mostConstrainedKeyword(state, alpha):
    # State spaces is a dictionary where the keys
    # are the keywords and the values are a list
    # of possible links
    stateSpaces = {}
    # We initialize the state spaces for the
    # unassigned keywords
    for keyword in unassignedKeywords(state):
        stateSpaces[keyword] = wikipedia.search(keyword)

    # Now for each keyword, we restrict its state
    # space based on all the other assigned keywords
    for keyword in stateSpaces:
        for assignedKeyword in assignedKeywords(state):
            # We determine how related the two words are
            # to each other and only restrict the keyword's
            # state space if the score is greater than
            # a predetermined threshold
```

```

score = linkingDependenceScore(keyword, assignedKeyword)
if score >= alpha:
    restrictStateSpace(keyword, assignedKeyword)

return keyword with min(stateSpaces)

```

2.3 The Relevance Function

Given a link, a keyword, and a current state, how do we determine if that link is the best fit for that keyword effectively using the context given to us? In our implementation, there is a crucial consideration: we must be able to somehow use previous assigned keywords as factors in our score. This means that given the content of a wikipedia link, we must use the current keyword and the keywords around it to gauge how good of a match that link is to our current keyword. To accomplish this task, we will discuss several implementations of different measures of semantic-relatedness and approaches that have been used in literature to take advantage of the structured corpora of wikipedia.

Formally, we define the relevance function as

$$R : \mathcal{S} \times \mathcal{L} \times \mathcal{W} \rightarrow [0, 1]$$

Where \mathcal{S} is the state space of our search formalism

\mathcal{L} is the space of candidate links a particular keyword can take

\mathcal{W} is the space of keywords in our given state, which are simply the keys of our dictionary in our implementation.

Our relevance function formally abstracts the notion of semantic relatedness by disambiguating a keyword and link and capturing the best match with a score assigned to it. To begin to implement it, we define the following key terms

$$\eta = \{w \in \mathcal{W} - \{c\} \mid c = \text{current keyword}\}$$

$$\text{Similarity} = \cos(A, B) = \frac{A \cdot B}{\|A\| \|B\|} = \frac{\sum_{i=1}^n A_i B_i}{\sqrt{\sum_{i=1}^n A_i^2} \cdot \sqrt{\sum_{i=1}^n B_i^2}}$$

Where η we define as the *context* of a given keyword and *Similarity* is the *cosine similarity* between two vectors, which is chosen to enforce a relevance score between 0 and 1. We will also use Term-Frequency Inverse-Document Frequency (TF-IDF) - a statistic that scores the importance of certain terms based on two metrics

1. weighing a term up if it appears frequently in a document
2. weighing a term down if it appears frequently in multiple documents. This will mean it is most likely not a unique keyword

$$\text{TF-IDF} = \sqrt{f(k, p)} \cdot \log \left(\frac{|\mathcal{D}|}{|\{d \in \mathcal{D} \mid k \in \mathcal{D}\}|} \right)$$

Where k is the current keyword, p is its candidate link's wikipedia page content, $f(k, p)$ is the term frequency of k in p , $|\mathcal{D}|$ is the total number of wikipedia pages of the other keywords in the state, and $\{d \in \mathcal{D} \mid k \in d\}$ is the number of wikipedia pages corresponding to each of the other keywords that contain k (Chisholm 149).

The key to an effective relevance function will involve taking into account the similarity of the context of a given keyword to that keywords wikipedia entry using cosine similarity and encoding terms with TF-IDF weights. There are variations of this approach that have been successful in literature, but one that stands out over others (Citation Needed. It is the algorithm comparison paper) in terms of time and space complexity is known as Explicit Semantic Analysis. Dojchinovski provides a summary of the method, which we will work into our formulation and modify to implement our final relevance function through the following high-level description:

A selected keyword is preprocessed and represented as a TF-IDF term vector. For each keyword $w_i \in \eta$, the method retrieves the possible Wikipedia articles from the Wikipedia API c_1, \dots, c_n containing w_i . The semantic relatedness of the word w_i with candidate link c_j is computed such that the similarity of association between w_i and c_j is multiplied with the TF-IDF weight of w_i . The relatedness score for any two documents is determined by computing the cosine similarity between the vectors of this candidate link and keyword semantic relatedness. (Dojchinovski 3).

This is how the relevance score will be assigned and calculated when each node is expanded into child nodes within the classical search formulation.

3. Algorithmic Analysis & Evaluation

As a way of assessing algorithmic efficiency, we begin by determining the time complexity of the approaches we followed. We know from AIMA that Depth First Search runs in polynomial time, specifically $O(b^d)$ where b is the branching factor and d is the depth of our solution. Our choice of encoding states as dictionaries of keyword-link mappings allows many of our subroutines to have constant time access to viewing and modifying these states, reducing the overhead from the problem formulation. However, there are additional factors in our implementation that can significantly slow down its run-time. The most-constrained keyword heuristic runs at $O(n + s \cdot n) = O(s \cdot n)$ complexity where $n = |\text{UNASSIGNEDKEYWORDS}|$, the length of the space of unassigned keywords and $s = |\text{ASSIGNEDKEYWORDS}|$ - an inefficient O .

Additionally, the relevance function subroutine utilizes TF-IDF in its implementation, an algorithm known to have a worst-case running time of $O(|N| \cdot |V|)$ where N is the corpus of the wikipedia page of a keyword and V is the space of all candidate links for that keyword. Using an inverted index, however, can improve TF-IDF to $O(|B| \cdot |V|)$ where B is the average over all terms in the query document of the number of other documents in which a keyword appears (Salakhutdinov 2). Reducing the overhead of all these subroutines will be implementation-dependent and will only be practically efficient for a set of input text with very small cardinality.

3.1 Implementation Critique

Our implementation is highly constrained by the rate-limit and speed of the API calls we use from the Wikipedia python library.

4. Algorithmic Limitations

Currently our search formulation is highly parameter-dependent. The accuracy of our entity-linking directly relies on the choice of γ (the relevance score threshold), which has to do with determining how high a value of a relevance score will determine a terminal state in the search problem, and ψ (the linking-dependence threshold) which determines how related two keywords are to each other in a given state.

Additionally, our system only works for single-word entity linking. This means that we cannot link bigrams such as “Steve Jobs” or any n-grams for $n > 1$. The algorithms are also limited to the extent of how powerful the wikipedia API is to fetch the set of candidate links it believes best match a given keyword and we depend on this for the expansion of a node within the different search algorithms.

5. System Feature Testing

Our system testing will rely on measuring the running time of multiple examples provided to the system using 4 key metrics as described in (Citation Needed). We can now frame our parameter-based system as an optimization problem, where we find the optimal values of γ and ψ that maximize candidate accuracy for values from 0 to 1 in our system. We run our system and measure these different metrics

6. Key Takeaways From Implementation

Takeaways

7. Empirical & Graphical Results

8. Appendix I

A trace of the program showing how it handles key examples or some other demonstration of the program in action.

9. Appendix II

A clear description of how to use your system and how to generate the output you discussed in the write-up and the example transcript in Appendix 1. N.B.: The teaching staff must be able to run your system.

10. Appendix III

A list of each project participant and that participants contributions to the project. If the division of work varies significantly from the project proposal, provide a brief explanation.